


Lecture 15: Optimization

CS 109B, STAT 121B, AC 209B, CSE 109B

Mark Glickman and Pavlos Protopapas



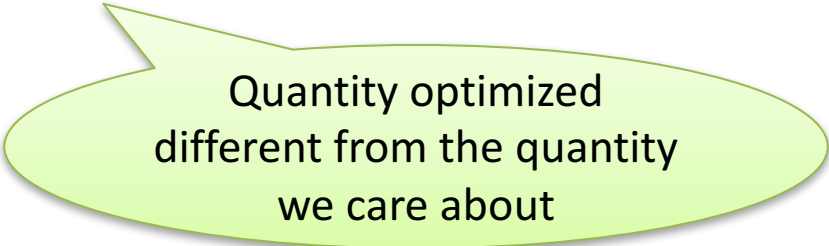
Learning vs. Optimization

- Goal of learning: minimize generalization error

$$J(\theta) = \mathbf{E}_{(x,y) \sim p_{data}} [L(f(x;\theta), y)]$$

- In practice, empirical risk minimization:

$$\hat{J}(\theta) = \frac{1}{m} \sum_{i=1}^m L(f(x^{(i)}; \theta), y^{(i)})$$



Quantity optimized
different from the quantity
we care about

Batch vs. Stochastic Algorithms

Shuffle the data before mini-batch

- Batch algorithms
 - Optimize empirical risk using **exact gradients**
- Stochastic algorithms
 - Estimates gradient from a **small random sample**

$$\nabla J(\theta) = \mathbf{E}_{(x,y) \sim p_{data}} [\nabla L(f(x;\theta), y)]$$


Large mini-batch: gradient computation expensive

Small mini-batch: greater variance in estimate,
longer steps for convergence

Critical Points

- Points with **zero gradient**
- 2nd-derivate (Hessian) determines curvature

- Saddle point and min and max
- Derivative may be 0 but point is max



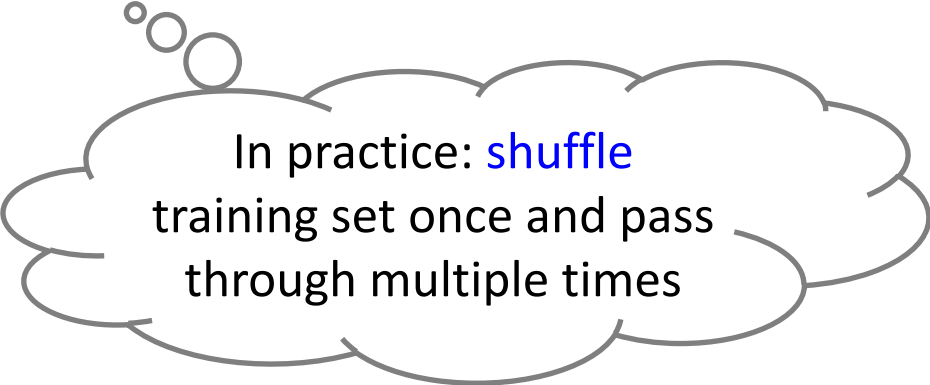
Stochastic Gradient Descent

- Take small steps in direction of **negative gradient**
- Sample m examples from training set and compute:

$$g = \frac{1}{m} \sum_i \nabla L(f(x^{(i)}; \theta), y^{(i)})$$


- Update parameters:

$$\theta = \theta - \varepsilon_k g$$



In practice: **shuffle**
training set once and pass
through multiple times

Stochastic Gradient Descent

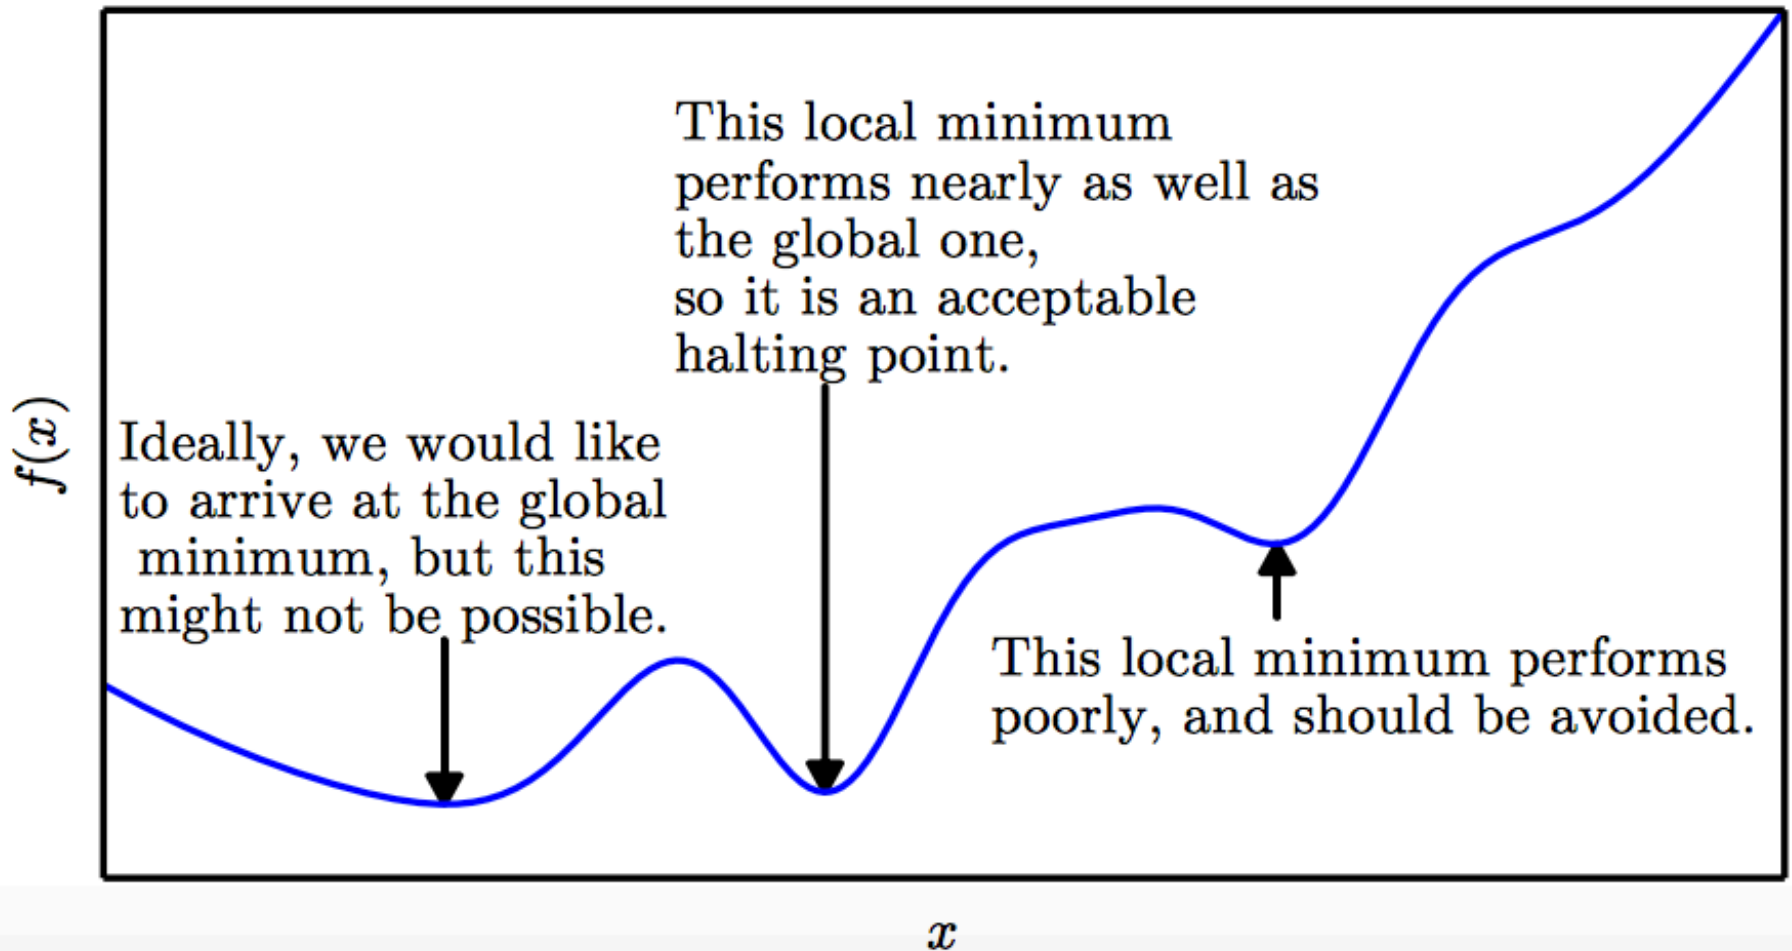


Oscillations because
updates do not exploit
curvature information

Outline

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization


Local Minima



Local Minima

- Old view: local minima is major problem in neural network training
- Recent view:
 - For sufficiently large neural networks, **most local minima incur low cost**
 - Not important to find true global minimum


Saddle Points



- Recent studies indicate that in high dim, saddle points are more likely than local min
- Gradient can be very small near saddle points

Saddle Points

- SGD is seen to escape saddle points
 - Moves down-hill, uses noisy gradients




- Second-order methods get stuck
 - solves for a point with zero gradient

Poor Conditioning


- Poorly conditioned Hessian matrix
 - High curvature: small steps leads to huge increase
- Learning is slow despite strong gradients

Oscillations slow
down progress



No Critical Points

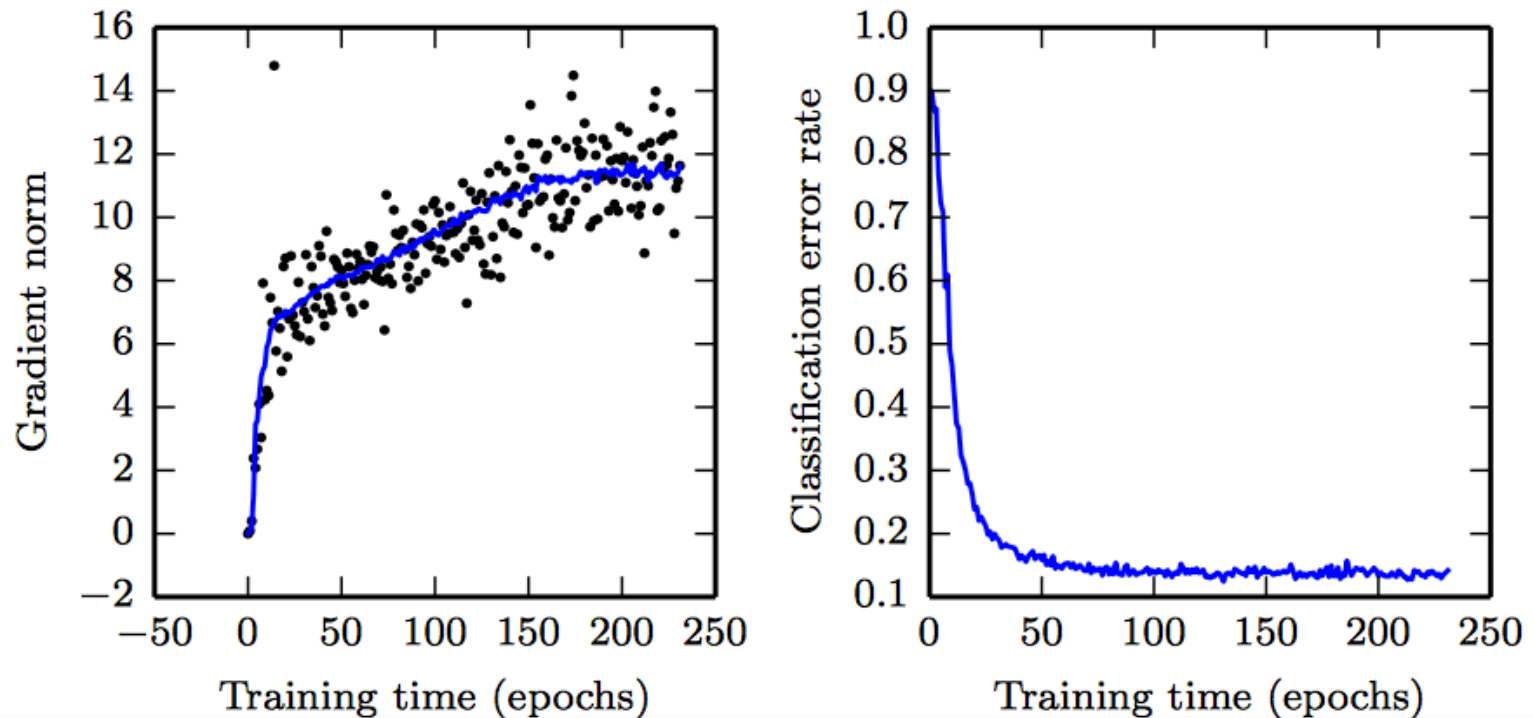
- Some cost functions do not have critical points



Early stopping works well

No Critical Points


Gradient norm increases, but validation error decreases



Convolution Nets for Object Detection

Goodfellow et al. (2016)

Exploding and Vanishing Gradients



Early stopping not useful if convergence before it

$$\mathbf{h}_1 = \mathbf{W}\mathbf{x}$$

$$\mathbf{h}_i = \mathbf{W}\mathbf{h}_{i-1}, \quad i = 2 \dots n$$

$$y = \sigma(h_1^n + h_2^n), \quad \text{where } \sigma(s) = \frac{1}{1 + e^{-s}}$$

Linear
activation

Exploding and Vanishing Gradients

Suppose $\mathbf{W} = \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$:

$$\begin{bmatrix} h_1^1 \\ h_2^1 \end{bmatrix} = \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \quad \dots \quad \begin{bmatrix} h_1^n \\ h_2^n \end{bmatrix} = \begin{bmatrix} a^n & 0 \\ 0 & b^n \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

$$y = \sigma(a^n x_1 + b^n x_2)$$

$$\nabla y = \sigma'(a^n x_1 + b^n x_2) \begin{bmatrix} na^{n-1} x_1 \\ nb^{n-1} x_2 \end{bmatrix}$$

Exploding and Vanishing Gradients

Suppose $x = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$

Case 1: $a = 1, b = 2$:


$$y \rightarrow 1, \quad \nabla y \rightarrow \begin{bmatrix} n \\ n2^{n-1} \end{bmatrix} \quad \text{Explodes!}$$

Case 2: $a = 0.5, b = 0.9$:

$$y \rightarrow 0, \quad \nabla y \rightarrow \begin{bmatrix} 0 \\ 0 \end{bmatrix} \quad \text{Vanishes!}$$


Exploding and Vanishing Gradients

- Exploding gradients lead to cliffs
- Can be mitigated using **gradient clipping**



Poor correspondence between local and global structure

It is really difficult to train large neural networks




Outline

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization

Momentum

- SGD is slow when there is **high curvature**



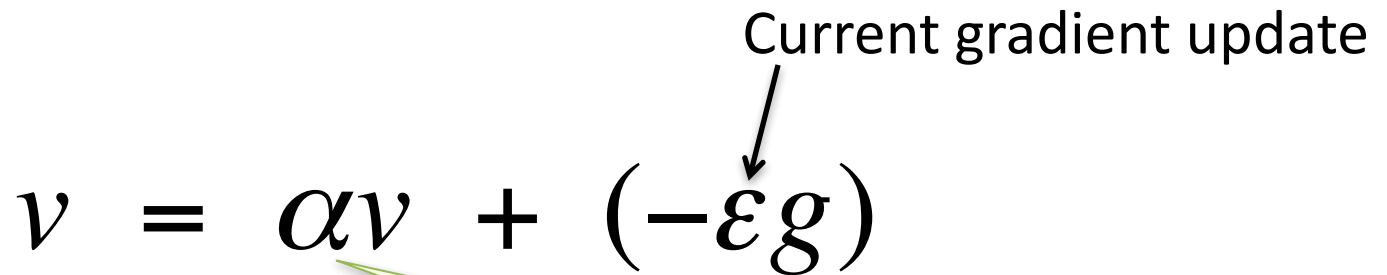
- Average gradient presents faster path to opt:
 - vertical components cancel out

Momentum

- Uses **past gradients** for update
- Maintains a new quantity: '**velocity**'
- *Exponentially decaying average* of gradients:

$$v = \alpha v + (-\varepsilon g)$$

Current gradient update



$\alpha \in [0,1)$ controls how quickly
effect of past gradients decay

Momentum

- Compute gradient estimate:

$$g = \frac{1}{m} \sum_i \nabla_{\theta} L(f(x^{(i)}; \theta), y^{(i)})$$

- Update velocity:

$$v = \alpha v - \epsilon g$$


- Update parameters:

$$\theta = \theta + v$$

Momentum

$J(\theta)$

Damped oscillations:
gradients in opposite
directions get
cancelled out



Nesterov Momentum

- Apply an **interim** update:

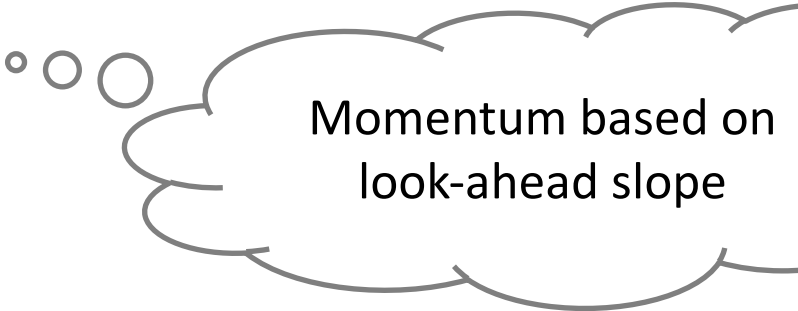
$$\tilde{\theta} = \theta + v$$

- Perform a correction based on gradient at the interim point:

$$g = \frac{1}{m} \sum_i \nabla_{\theta} L(f(x^{(i)}; \tilde{\theta}), y^{(i)})$$

$$v = \alpha v - \epsilon g$$

$$\theta = \theta + v$$




Momentum based on
look-ahead slope

Outline

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization

Adaptive Learning Rates

$J(\theta)$



- Oscillations along vertical direction
 - Learning must be slower along parameter 2
 - Use a different learning rate for each parameter?
- Derivative v. steep,
- If learning rate too large, bounces around

AdaGrad

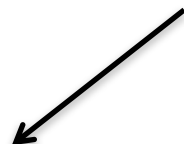
- Accumulate squared gradients:

$$r_i = r_i + g_i^2$$

- Update each parameter:

$$\theta_i = \theta_i - \frac{\varepsilon}{\delta + \sqrt{r_i}} g_i$$

Inversely
proportional to
cumulative
squared gradient



- Greater progress along gently sloped directions

RMSProp

- For non-convex problems, AdaGrad can prematurely decrease learning rate
- Use **exponentially weighted average** for gradient accumulation

$$r_i = \rho r_i + (1 - \rho) g_i^2$$

$$\theta_i = \theta_i - \frac{\varepsilon}{\delta + \sqrt{r_i}} g_i$$

Adam

- RMSProp + Momentum

- Estimate first moment:

$$v_i = \rho_1 v_i + (1 - \rho_1) g_i$$

- Estimate second moment:

$$r_i = \rho_2 r_i + (1 - \rho_2) g_i^2$$

- Update parameters:

$$\theta_i = \theta_i - \frac{\varepsilon}{\delta + \sqrt{r_i}} v_i$$

Also applies
bias correction
to v and r

Works well in practice,
is fairly robust to
hyper-parameters

Outline

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization

Parameter Initialization

- Goal: **break symmetry** between units
 - so that each unit computes a different function
- Initialize all weights (not biases) **randomly**
 - Gaussian or uniform distribution
- **Scale of initialization?**
 - *Large* -> grad explosion, *Small* -> grad vanishing

Xavier Initialization

- Heuristic for all outputs to have **unit variance**
- For a fully-connected layer with m inputs:

$$W_{ij} \sim N\left(0, \frac{1}{m}\right)$$

- For ReLU units, it is recommended:

$$W_{ij} \sim N\left(0, \frac{2}{m}\right)$$

Normalized Initialization

- Fully-connected layer with m inputs, n outputs:

$$W_{ij} \sim U\left(-\sqrt{\frac{6}{m+n}}, \sqrt{\frac{6}{m+n}}\right)$$

- Heuristic trades off between initialize all layers have same activation and gradient variance
- **Sparse** variant when m is large
 - Initialize k nonzero weights in each unit

Bias Initialization

- Output unit bias
 - Marginal statistics of the output in the training set
- Hidden unit bias
 - Avoid saturation at initialization
 - E.g. in ReLU, initialize bias to 0.1 instead of 0
- Units controlling participation of other units
 - Set bias to allow participation at initialization

- If initialised with only $W=0$, we want biases to be as close to std dev, mean as possible.

Outline

- Challenges in Optimization
- Momentum
- Adaptive Learning Rate
- Parameter Initialization
- Batch Normalization

Feature Normalization

- Good practice to normalize features before applying learning algorithm:

- NW: all nodes interconnected


$$x' = \frac{x - \mu}{\sigma}$$

Feature vector $\rightarrow x$ Vector of mean feature values $\leftarrow \mu$


σ Vector of SD of feature values $\leftarrow \sigma$

- Features in **same scale**: mean 0 and variance 1
 - Speeds up learning

Feature Normalization




Before normalization



After normalization

Internal Covariance Shift

Each hidden layer changes distribution of inputs to next layer: *slows down learning*



Batch Normalization

- Training time:
 - Mini-batch of activations for layer to normalize

$$H = \begin{bmatrix} H_{11} & \cdots & H_{1K} \\ \vdots & \ddots & \vdots \\ H_{N1} & \cdots & H_{NK} \end{bmatrix}$$

K hidden layer activations

N data points in mini-batch

Batch Normalization

- Training time:
 - Mini-batch of activations for layer to normalize

= Loss of expressiveness with normalization

$$H' = \frac{H - \mu}{\sigma}$$

where

$$\mu = \frac{1}{m} \sum_i H_{i,:}$$

Vector of mean activations
across mini-batch

$$\sigma = \sqrt{\frac{1}{m} \sum_i (H - \mu)_i^2 + \delta}$$

Vector of SD of each unit
across mini-batch

Batch Normalization

- Training time:
 - Normalization can reduce expressive power
 - Instead use:

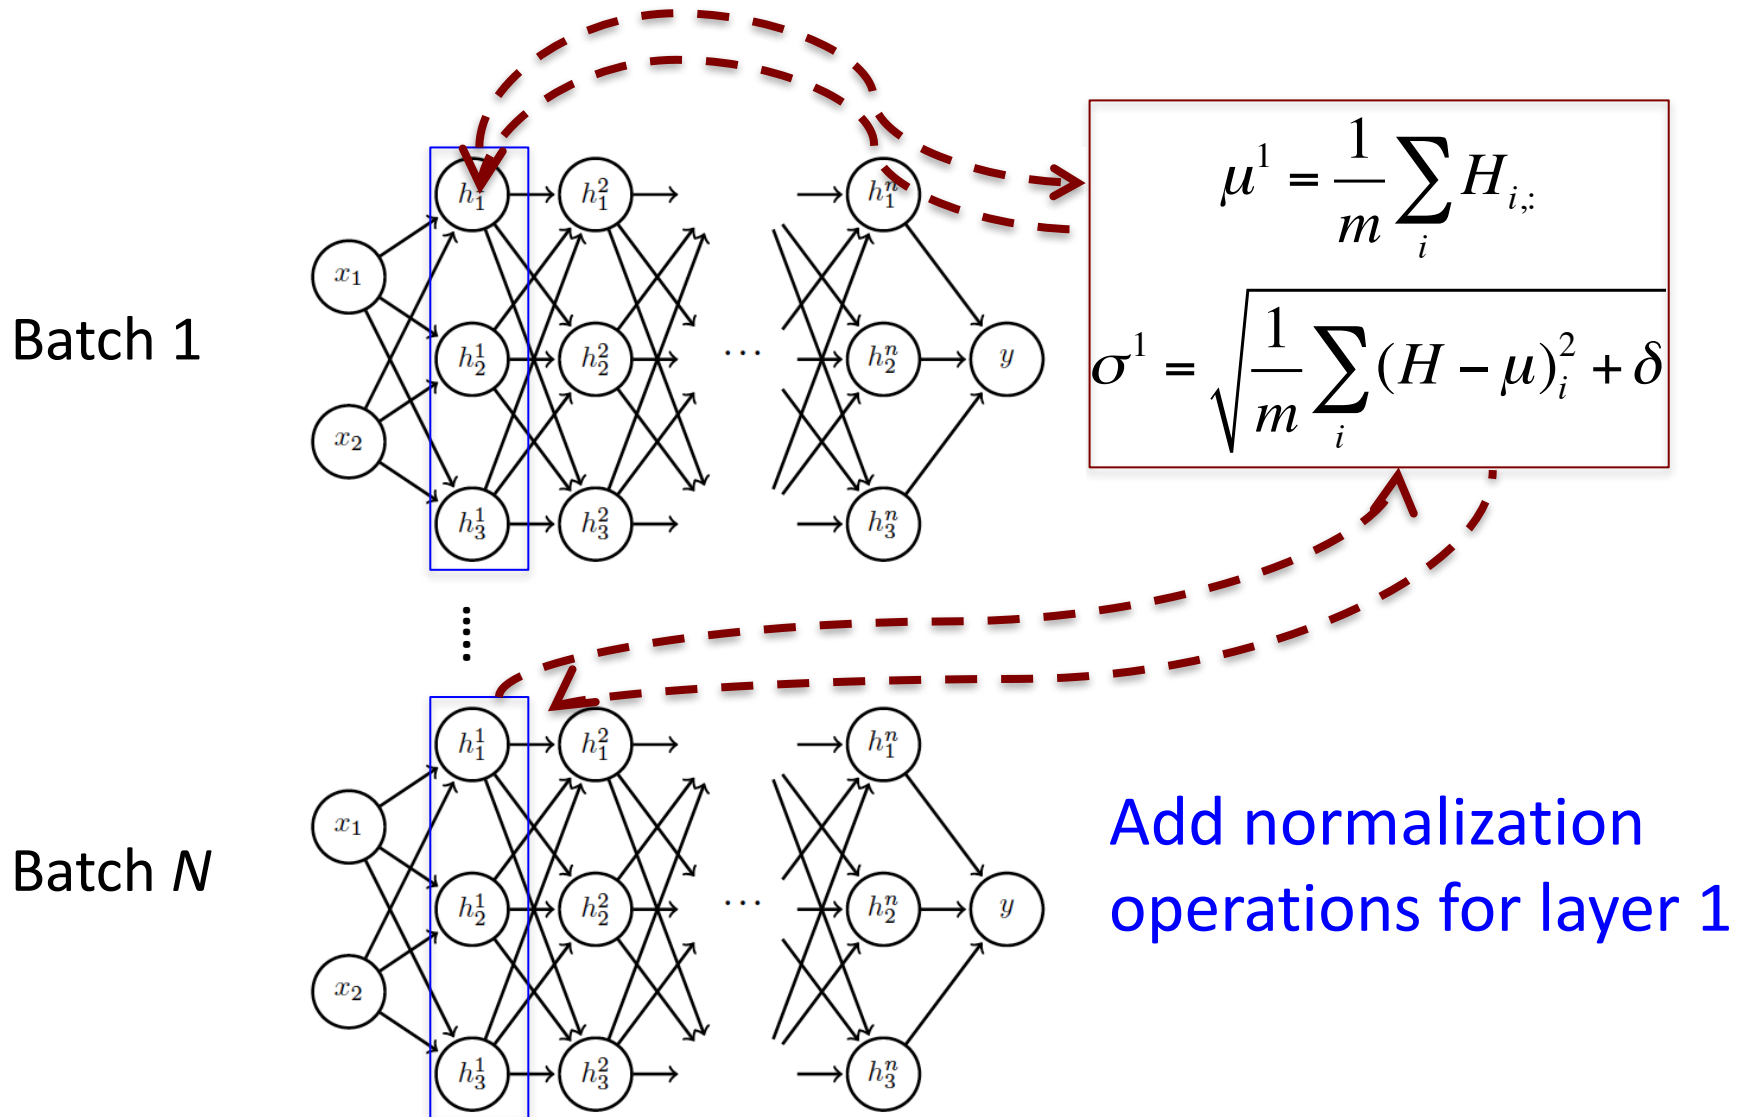
$$\gamma H' + \beta$$

- We get mu and sigma back.

Learnable parameters


- Allows network to **control range of normalization**

Batch Normalization




Batch Normalization

Batch 1



$$\mu^2 = \frac{1}{m} \sum_i H_{i,:}$$
$$\sigma^2 = \sqrt{\frac{1}{m} \sum_i (H - \mu)_i^2 + \delta}$$

Batch N



Add normalization
operations for layer 2
and so on ...

Batch Normalization

- Differentiate the **joint loss** for N mini-batches
- Back-propagate *through* the norm operations
- Test time:
 - Model needs to be evaluated on a *single example*
 - Replace μ and σ with **running averages** collected during training